

Let's Try Some Bandits (practical)

Tom Vodopivec

IADS Big Data and Analytics Summer School
2019-08-05

////////



Overview

We have a set multi-armed bandits inspired by real data (obtained from our Dynamic Creative Optimization system).

Stochastic with reward either 0 or 1 (Bernoulli bandit)

Budget-limited (from 500 up to 30000 pulls)

10 stationary cases

10 non-stationary cases (multiple segments with constant probability)

Practical Work

1. [Download the code](#), test it executes successfully
2. Implement and measure the performance of bandit policies:
 - a. Epsilon-greedy
 - b. UCB1
 - c. Thompson sampling
 - d. (advanced) UCB-Tuned
3. Test different algorithms' parameters
4. Own ideas for further improvement?

Use [Main.py](#) to run and configure your experiment

Add your bandit policies to [Agent.py](#)